


Univariate Linear Regression

Regression and machine learning

- Regression is a process of estimating the relationship between a dependent variable (or outcome variable) and independent variables (or features, predictors...)
- Linear regression – the dependent variable relates to the independent variables linearly.
- In the context of machine learning, regression is used for prediction/forecasting.

Housing Prices (Portland, OR)



Supervised learning

Given data pairs (size, price),
create parameterized models by
supervised learning

Simplest such model,
univariate linear regression:

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Regression problem

Predict real valued output (price)
given the size of a house just put
on market

Classification problem

Predict discrete valued output (top
moon/bottom moon) given the
coordinates of a data point

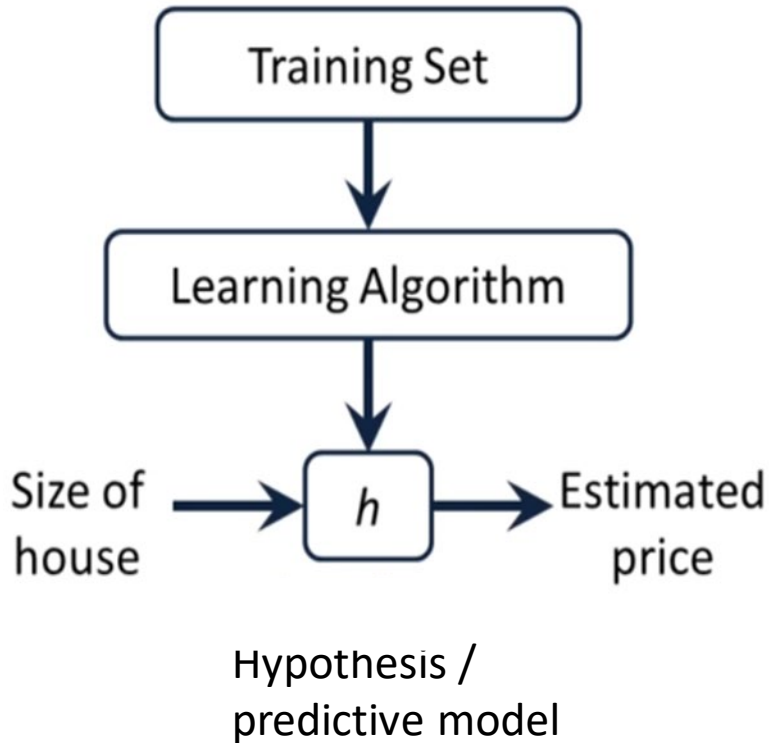
Training set of housing prices (Portland, OR)	Size in feet ² (x)	Price (\$) in 1000's (y)
	2104	460
	1416	232
	1534	315
	852	178

Notation:

m: number of training examples

x: input variable or input feature or feature

y: output variable



To represent h – need a structure/model & model parameters

Univariate linear regression:

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

x : size of house

θ : parameter of model

How to choose model parameters θ_0 and θ_1 ?


Goal: to make $h_{\theta}(x)$ as close to real house price (y) as possible (using the training data)

Problem formulation by first defining a cost function, $J(\theta_0, \theta_1)$:

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Choose θ_0 and θ_1 so that is $J(\theta_0, \theta_1)$ minimized

Let's use examples to gain some intuition about the cost function $J(\theta_0, \theta_1)$
As a function of the parameters θ_0 and θ_1



Hypothesis: $h_{\theta}(x) = \theta_0 + \theta_1 x$

Parameters: θ_0, θ_1

Cost function: $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$

Goal: minimize $J(\theta_0, \theta_1)$
 θ_0, θ_1

Outline of learning procedure:

- Start with some θ_0, θ_1
- Keep changing θ_0, θ_1 to reduce $J(\theta_0, \theta_1)$ until we hopefully end up at a minimum

Update parameters θ_0 and θ_1 using gradient descent until convergence

Simultaneously update θ_0 and θ_1 according to

$$\theta_0 := \theta_0 - \eta \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1)$$
$$\theta_1 := \theta_1 - \eta \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1)$$

η is a positive number, which is called the learning rate
– too small, slow learning; too large, divergence

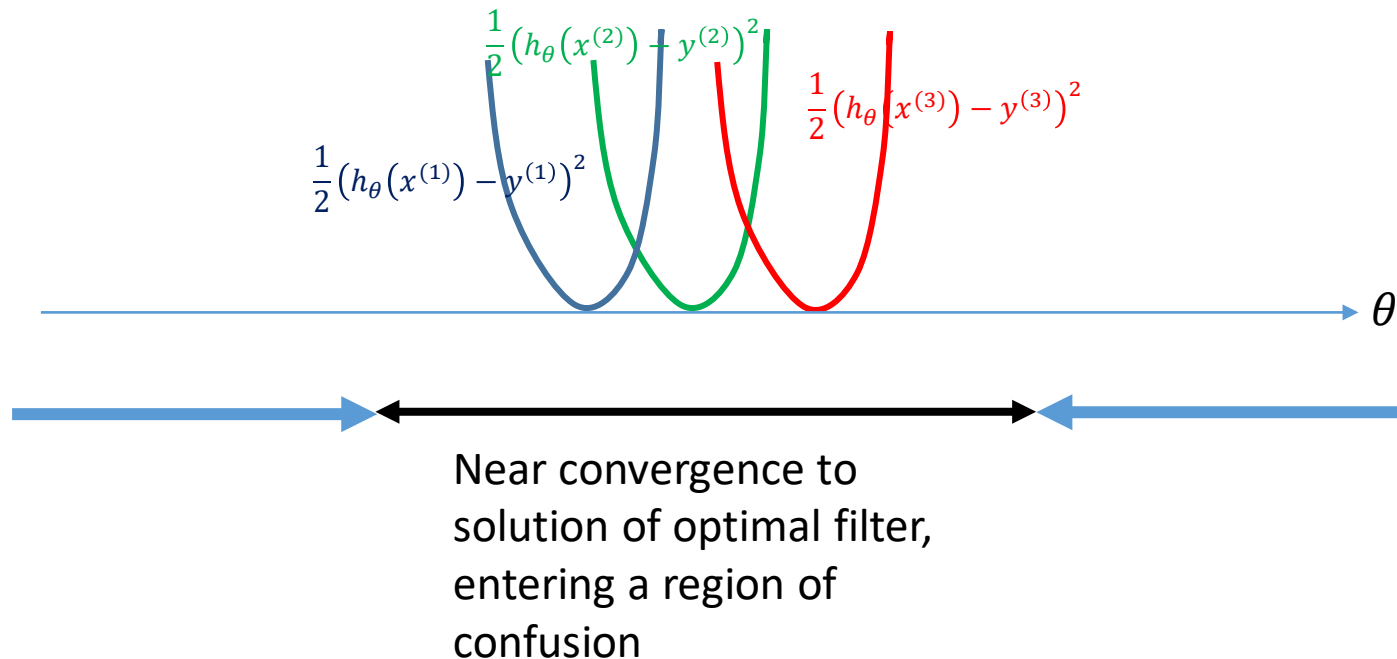
Exercise – derive the two gradients

- notice this derivation is a “batch” learning algorithm
- “online” learning – the LMS algorithm

How to optimize (minimize) a cost/loss function by tuning neural network weights for a predicted value to approach an actual value?

- Stochastic gradient descent (such as the LMS by Widrow and Hoff)
- Batch gradient descent as in linear regression, quadratic problem with least squares solution in closed form (may not be the best for large data set)
- Mini-batches (as in many deep networks)

Insights on SGD (using Widrow-Hoff LMS setting)



- When far from convergence, behaves similarly as ordinary GD but less computation ($1/m$, m is the # of samples)
- When near convergence, gradient decent for one learning sample may be an ascent for another sample. A diminishing step size is required, e.g., the Robbins-Monro condition (e.g., $\eta=1/n$, n is # of learning iteration)